

IMDB NLP

Jiaqi Sun, Huifei Xu, Hao He, Yaquan Yang

2022-10-16

In this topic modeling assignment, we use the IMDB Movie Reviews (50K) dataset from Kaggle to analyze the review sentiment and pull out independent topic categories(i.e., movie genre) based on group of keywords.

1. The tidy text format

First, we import data and break down each review into a list of words. Here, we tokenize at word level and treat each review as a separate “document”. We manually created a review number for each document since the movie titles for each review are not included in this dataset. Capitalization in reviews is kept for the initial screening.

```
IMDB <- read.csv("IMDB Dataset.csv")

IMDB_df <- tibble(IMDB)
# glimpse(IMDB_df)
# two columns: `sentiment` indicates whether the review is labeled negative or
# positive; `review` includes the text of each review.
head(IMDB_df)
```

```
## # A tibble: 6 × 2
##   review                                                                 senti...1
##   <chr>                                                                <chr>
## 1 "One of the other reviewers has mentioned that after watching just 1 ... positi...
## 2 "A wonderful little production. <br /><br />The filming technique is ... positi...
## 3 "I thought this was a wonderful way to spend time on a too hot summer... positi...
## 4 "Basically there's a family where a little boy (Jake) thinks there's ... negati...
## 5 "Petter Mattei's \"Love in the Time of Money\" is a visually stunning... positi...
## 6 "Probably my all-time favorite movie, a story of selflessness, sacrific... positi...
## # ... with abbreviated variable name `sentiment`
```

```
# create review id
IMDB_df %>%
  mutate(review_number = row_number()) ->IMDB_df
```

```

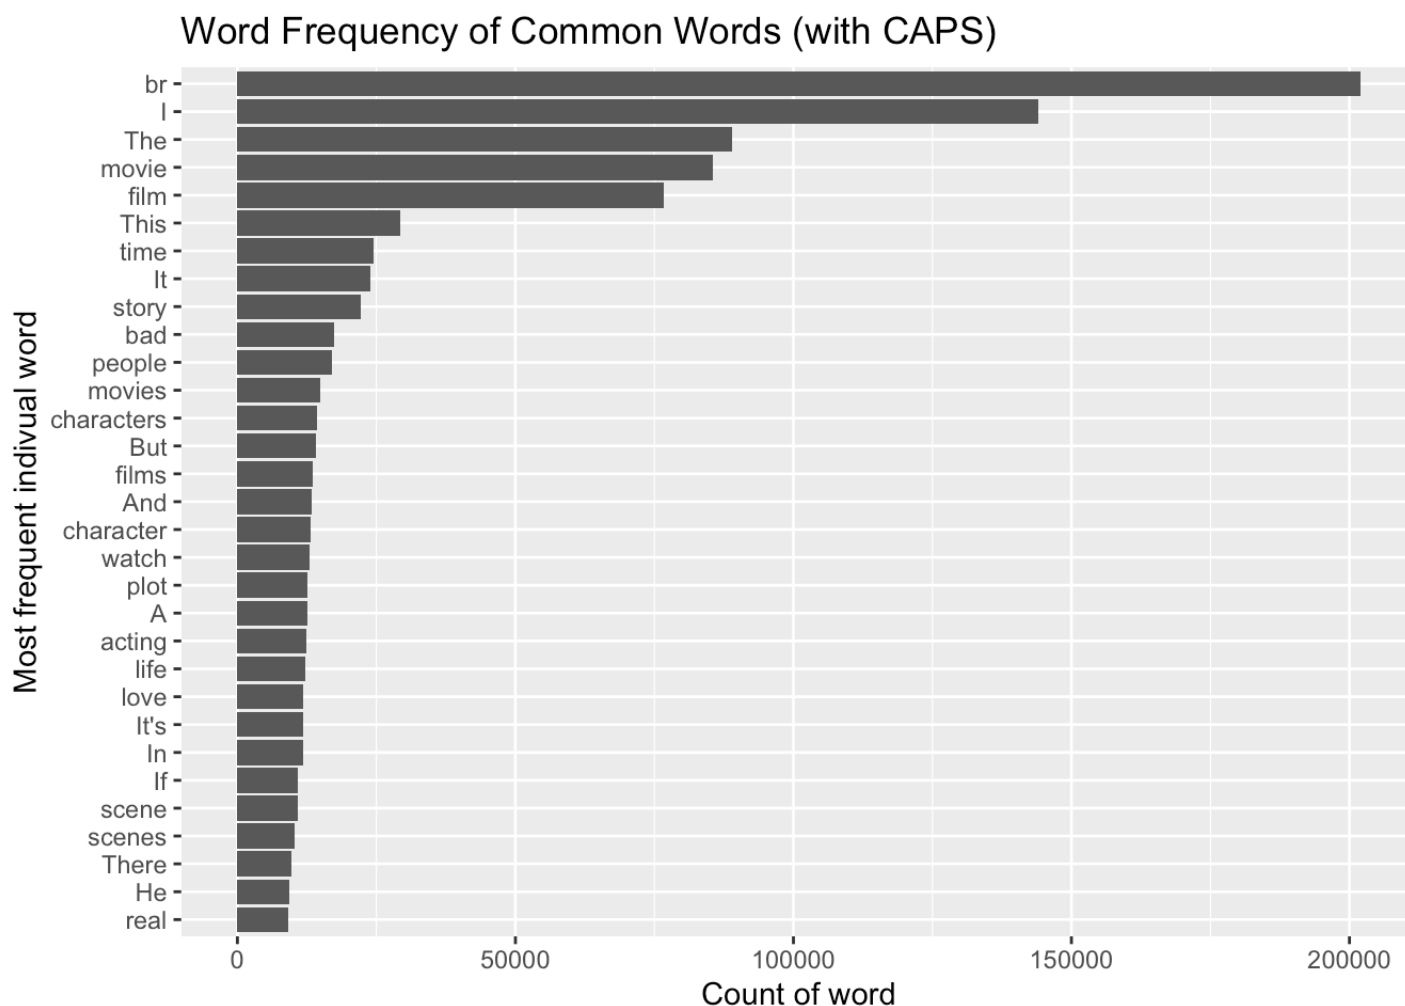
# keep the CAPS for sentiment analysis
tidy_IMDB <- IMDB_df %>%
  unnest_tokens(word, review, to_lower = FALSE)%>% anti_join(stop_words)

# tidy_IMDB

# most common words
# tidy_IMDB %>%
# count(word, sort = TRUE)

#plot frequencies
tidy_IMDB %>%
  count(word, sort = TRUE) %>%
  filter(n > 9000) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word)) +
  geom_col() +
  labs(y = NULL,
       title = "Word Frequency of Common Words (with CAPS)" +
  xlab("Count of word")+
  ylab("Most frequent indivual word")

```



From the frequency plot above, we can see that if we keep CAPS in each review, the most common words would include some words (such as pronouns) that are useless for sentiment analysis and topic identification, so we decide to lowercase all words.

```
tidy_IMDB <- IMDB_df %>%
  unnest_tokens(word, review)%>%
  anti_join(stop_words)

head(tidy_IMDB)
```

```
## # A tibble: 6 × 3
##   sentiment review_number word
##   <chr>          <int> <chr>
## 1 positive             1 reviewers
## 2 positive             1 mentioned
## 3 positive             1 watching
## 4 positive             1 1
## 5 positive             1 oz
## 6 positive             1 episode
```

`stop_words` is a data frame from `tidytext` package that contains English stop words from three lexicons. Except for these stop-words, we also customize a list of domain specific stop-words.

See below:

Add “br” to be a customized stop-word as it’s a leftover from html format. Add “film”, “movie”, “em” to customized stop-word as they are appearing in any review and meaningless

```
data(stop_words)
stop_words <- bind_rows(tibble(word = c("br", "film", "movie", "films",
"movies", "characters", "character", "story", "time", "people", "watching",
"scene", "scenes", "plot", "watch", "real", "cast", "director", "lot", "pretty",
"10", "actors", "1", "oz", "makes", "2", "em"), lexicon = c("custom")),
stop_words)

tidy_IMDB <- tidy_IMDB %>%
  anti_join(stop_words)

head(tidy_IMDB)
```

```
## # A tibble: 6 × 3
##   sentiment review_number word
##   <chr>           <int> <chr>
## 1 positive             1 reviewers
## 2 positive             1 mentioned
## 3 positive             1 episode
## 4 positive             1 hooked
## 5 positive             1 happened
## 6 positive             1 struck
```

After remove all the stop-words, we try stemming and lemmatization. Stemming using rules to cut words down to their stems, operating on the word by itself. Lemmatization uses knowledge about a language's structure to reduce words down to their lemmas, the canonical or dictionary forms of words, operating on the word in its context. As an important part of NLP pipelines, these methods would help us reduce the feature space of text data and should decrease the sparsity of text data. In this way, we may have a better fitted LDA model in the later process.

```
# stemming and lemmatization
library(textstem)
lemmatization <- tidy_IMDB %>% mutate(lemma = word%>% lemmatize_words())
lemmatization %>% head()
```

```
## # A tibble: 6 × 4
##   sentiment review_number word      lemma
##   <chr>           <int> <chr>    <chr>
## 1 positive             1 reviewers reviewer
## 2 positive             1 mentioned mention
## 3 positive             1 episode  episode
## 4 positive             1 hooked   hook
## 5 positive             1 happened happen
## 6 positive             1 struck   strike
```

```
library(SnowballC)
stemming <- tidy_IMDB %>% mutate(stem = wordStem(word))
stemming %>% head()
```

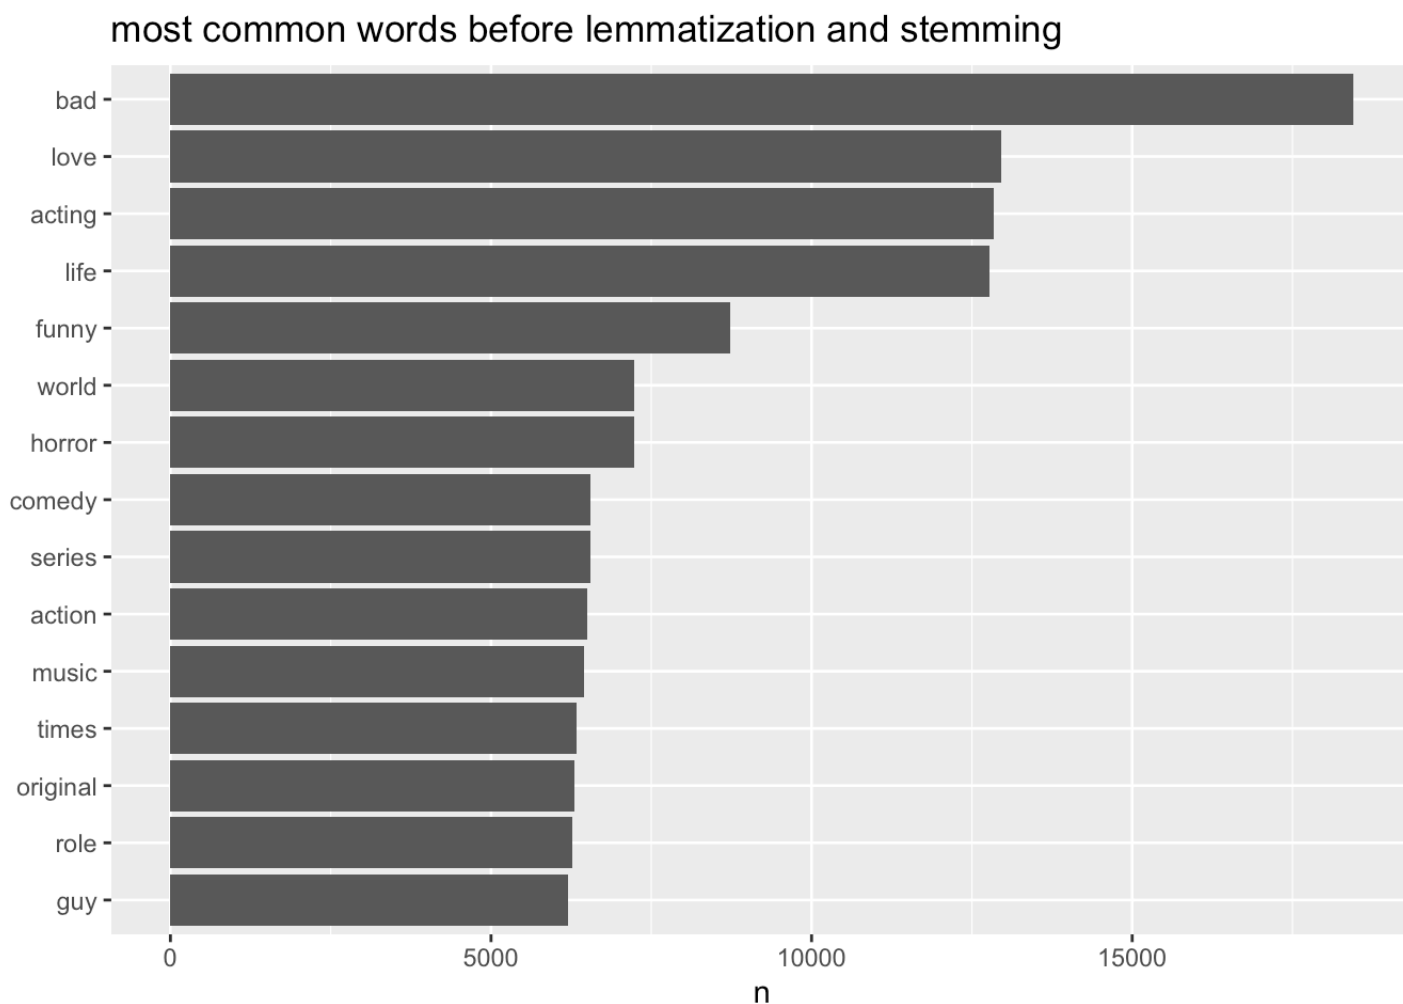
```
## # A tibble: 6 × 4
##   sentiment review_number word      stem
##   <chr>           <int> <chr>    <chr>
## 1 positive             1 reviewers review
## 2 positive             1 mentioned mention
## 3 positive             1 episode  episod
## 4 positive             1 hooked   hook
## 5 positive             1 happened happen
## 6 positive             1 struck   struck
```

```
tidy_IMDB %>%
  count(sentiment, word, sort = TRUE) %>% head()
```

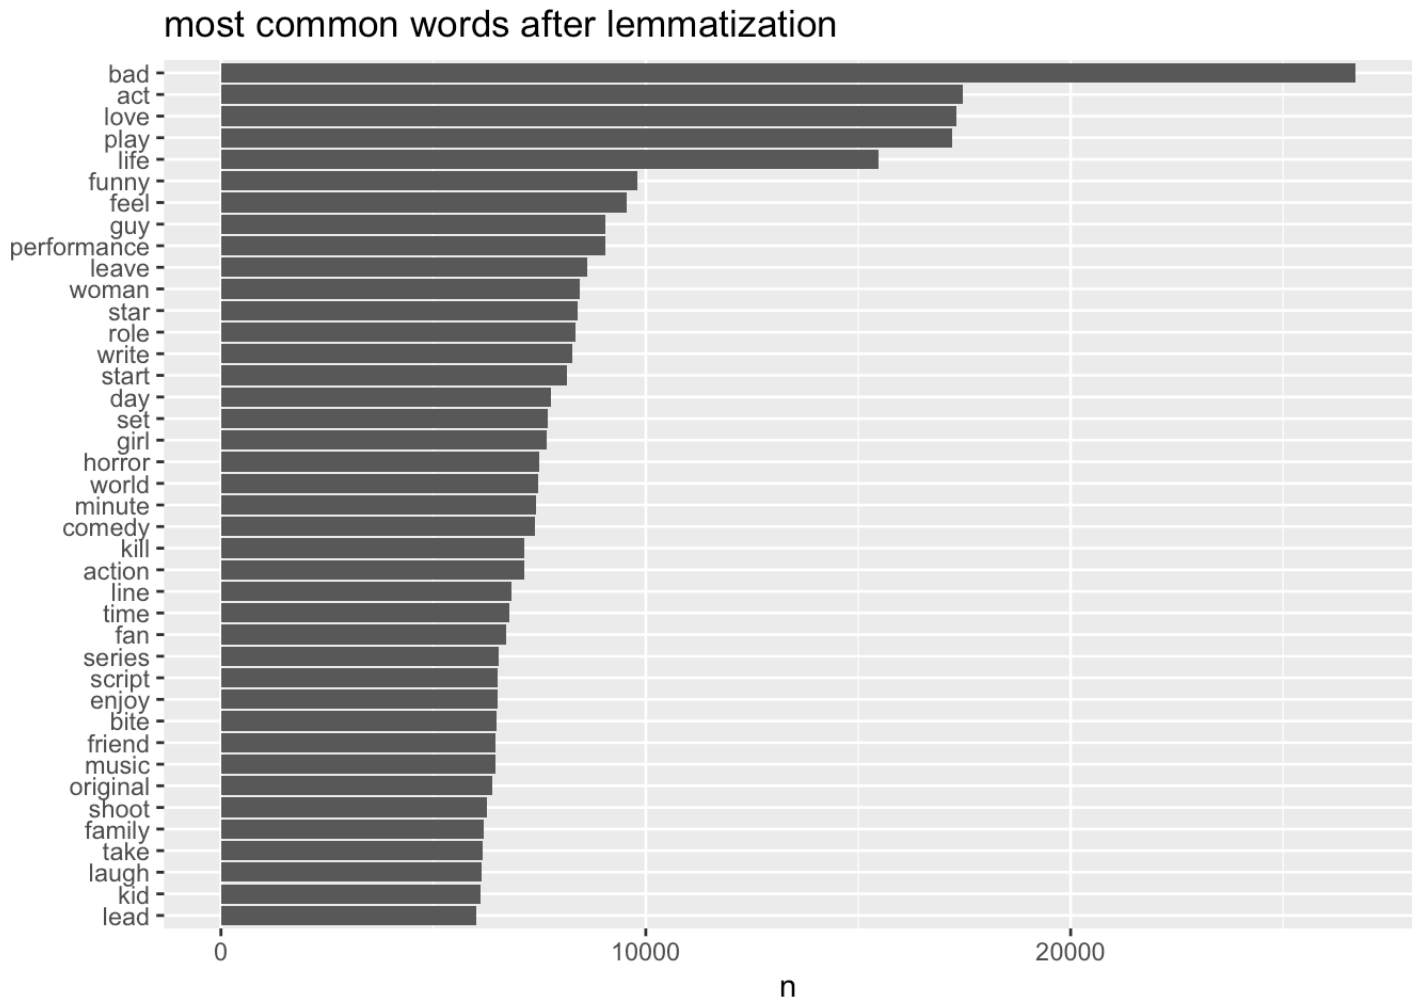
```
## # A tibble: 6 × 3
##   sentiment word      n
##   <chr>      <chr> <int>
## 1 negative  bad     14706
## 2 positive  love     8669
## 3 negative  acting   8070
## 4 positive  life     8040
## 5 negative  worst    4884
## 6 positive  acting   4772
```

#The most common words in IMDB comments

```
tidy_IMDB %>%
  count(word, sort = TRUE) %>%
  filter(n > 6000) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word)) +
  geom_col() +
  labs(y = NULL, title = "most common words before lemmatization and stemming")
```




```
# The most common words after lemmatization
lemmatization %>%
  count(lemma, sort = TRUE) %>%
  filter(n > 6000) %>%
  mutate(lemma = reorder(lemma, n)) %>%
  ggplot(aes(n, lemma)) +
  geom_col() +
  labs(y = NULL, title = "most common words after lemmatization")
```



```
# The most common words after stemming
stemming %>%
  count(stem, sort = TRUE) %>%
  filter(n > 8000) %>%
  mutate(stem = reorder(stem, n)) %>%
  ggplot(aes(n, stem)) +
  geom_col() +
  labs(y = NULL, title = "most common words after stemming")
```

most common words after stemming



It's easy to see that stemming would convert "plays" to "plai", while in lemmatization it would be "play".

```
tidy_IMDB %>%  
  count(sentiment)
```

```
## # A tibble: 2 × 2  
##   sentiment      n  
##   <chr>      <int>  
## 1 negative 1928848  
## 2 positive 2049952
```

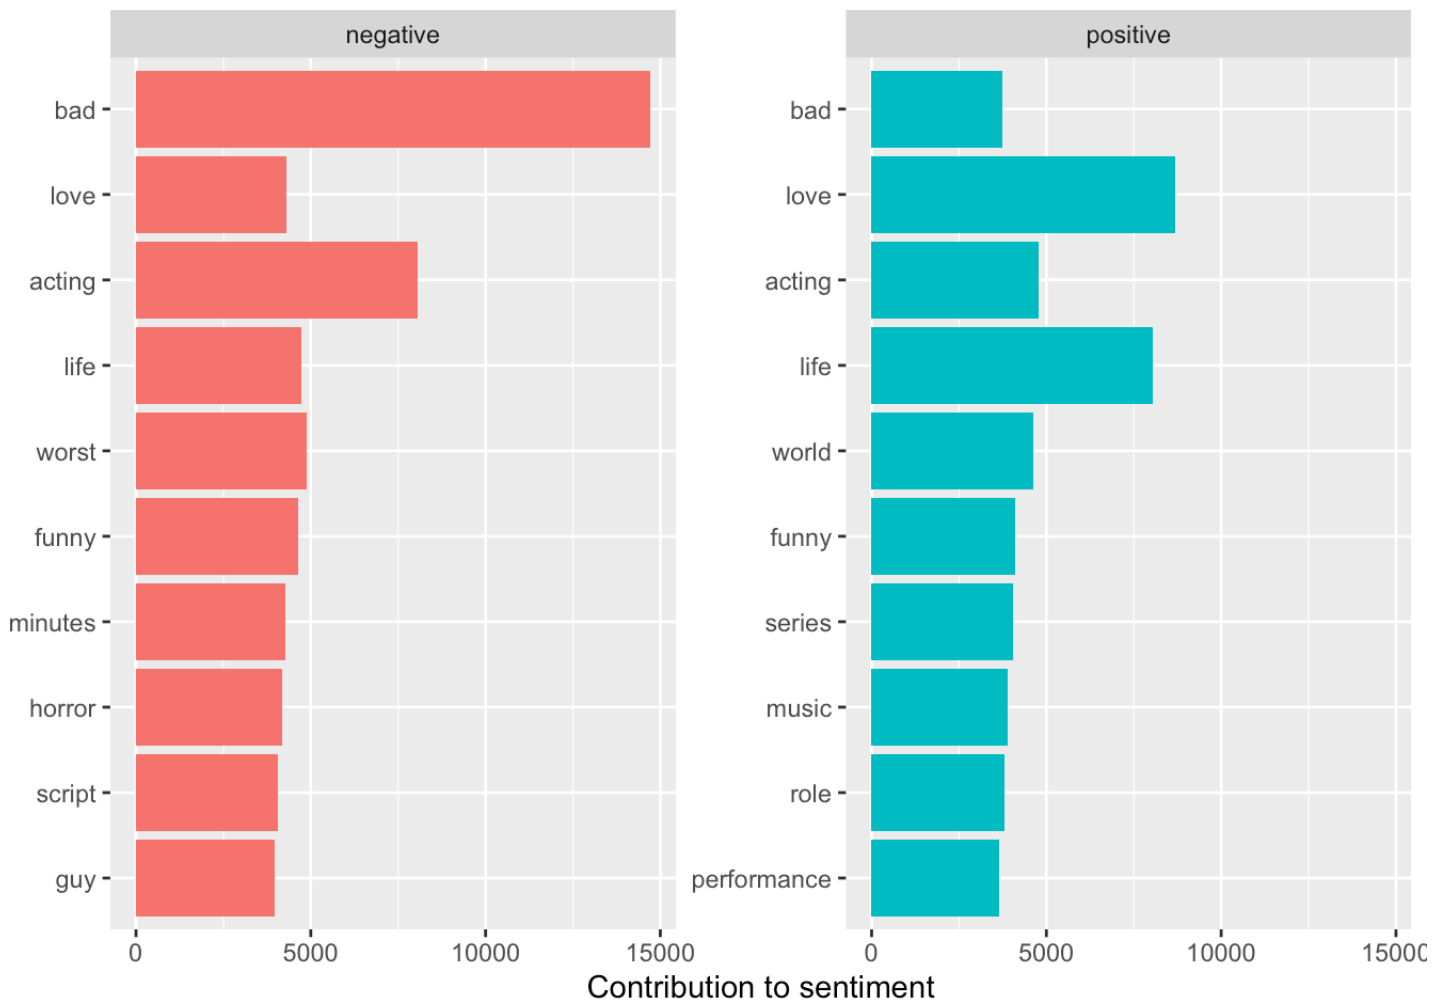
2. Sentiment analysis with tidy data

```
bing_word_counts <- tidy_IMDB %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()

bing_word_counts %>% head()
```

```
## # A tibble: 6 × 3
##   word      sentiment      n
##   <chr>    <chr>      <int>
## 1 bad      negative    14706
## 2 love     positive     8669
## 3 acting   negative     8070
## 4 life     positive     8040
## 5 worst    negative     4884
## 6 acting   positive     4772
```

```
#Words that contribute to positive and negative sentiment in IMDB
bing_word_counts %>%
  group_by(sentiment) %>%
  slice_max(n, n = 10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(x = "Contribution to sentiment",
       y = NULL)
```

Most common positive and negative words in IMDB

```
tidy_IMDB %>%  
  count(word, sentiment, sort = TRUE) %>%  
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%  
  comparison.cloud(colors = c("gray20", "gray80"),  
    max.words = 100)
```

negative



positive

```
# table(IMDB$sentiment)
```

The polarity function scans the subjectivity lexicon for positive and negative words, and we can further visualize the following results to compare with our sentiment tags. The subjective dictionary used by the qdap package contains about 6800 words with tags, which would be introduced here to help measuring the density of keywords

```
## [1] "C/C/C/C/C/en_US.UTF-8"
```

```
## Warning in polarity(as.character(IMDB$review)):
```

```
## Some rows contain double punctuation. Suggested use of `sentSplit` function.
```


```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
```

```
## i Please use `linewidth` instead.
```

```
## Warning: The dot-dot notation (`..density..`) was deprecated in ggplot2 3.4.0.
```

```
## i Please use `after_stat(density)` instead.
```

qdap's positive and negative word ratings



3. Analyzing word and document frequency: tf-idf

```
IMDB_tf_idf <- tidy_IMDB %>%  
  count(review_number, word, sort = TRUE) %>%  
  bind_tf_idf(word, review_number, n) %>%  
  arrange((tf_idf))
```

```
IMDB_tf_idf %>% head()
```

```
## # A tibble: 6 x 6  
##   review_number word      n      tf    idf  tf_idf  
##   <int> <chr>   <int>   <dbl> <dbl>   <dbl>  
## 1     31482 world      1 0.000899 2.22 0.00200  
## 2     31437 love       1 0.00117 1.73 0.00202  
## 3     31241 love       1 0.00118 1.73 0.00204  
## 4     40522 acting    1 0.00140 1.54 0.00217  
## 5     31482 tv        1 0.000899 2.47 0.00222  
## 6     31482 watched   1 0.000899 2.53 0.00227
```

```
# summary(IMDB_tf_idf)
```

The inverse document frequency (and thus tf-idf) is very low (near zero) for words that occur in many of the reviews in a collection; this is how this approach decreases the weight for common words. The inverse document frequency will be a higher number for words that occur in fewer of the documents in the collection.

4. Relationships between words: n-grams and correlations

```
IMDB_bigrams <- IMDB_df %>%  
  unnest_tokens(bigram, review, token = "ngrams", n = 2) %>%  
  filter(!is.na(bigram))  
  
IMDB_bigrams %>% head()
```

```
## # A tibble: 6 x 3  
##   sentiment review_number bigram  
##   <chr>          <int> <chr>  
## 1 positive          1 one of  
## 2 positive          1 of the  
## 3 positive          1 the other  
## 4 positive          1 other reviewers  
## 5 positive          1 reviewers has  
## 6 positive          1 has mentioned
```

```
IMDB_bigrams %>%  
  count(bigram, sort = TRUE) %>% head
```

```
## # A tibble: 6 x 2  
##   bigram      n  
##   <chr>    <int>  
## 1 br br    101039  
## 2 of the    77235  
## 3 in the    50216  
## 4 this movie 31166  
## 5 and the    26639  
## 6 is a      26080
```

As one might expect, a lot of the most common bigrams are pairs of common (uninteresting) words, such as of the and to be: what we call “stop-words” (see Chapter 1). This is a useful time to use tidyr’s separate(), which splits a column into multiple based on a delimiter. This lets us separate it into two columns, “word1” and “word2”, at which point we can remove cases where either is a stop-word.

```
bigrams_separated <- IMDB_bigrams %>%
  separate(bigram, c("word1", "word2"), sep = " ")

bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word)

# new bigram counts:
bigram_counts <- bigrams_filtered %>%
  count(word1, word2, sort = TRUE)

bigram_counts
```

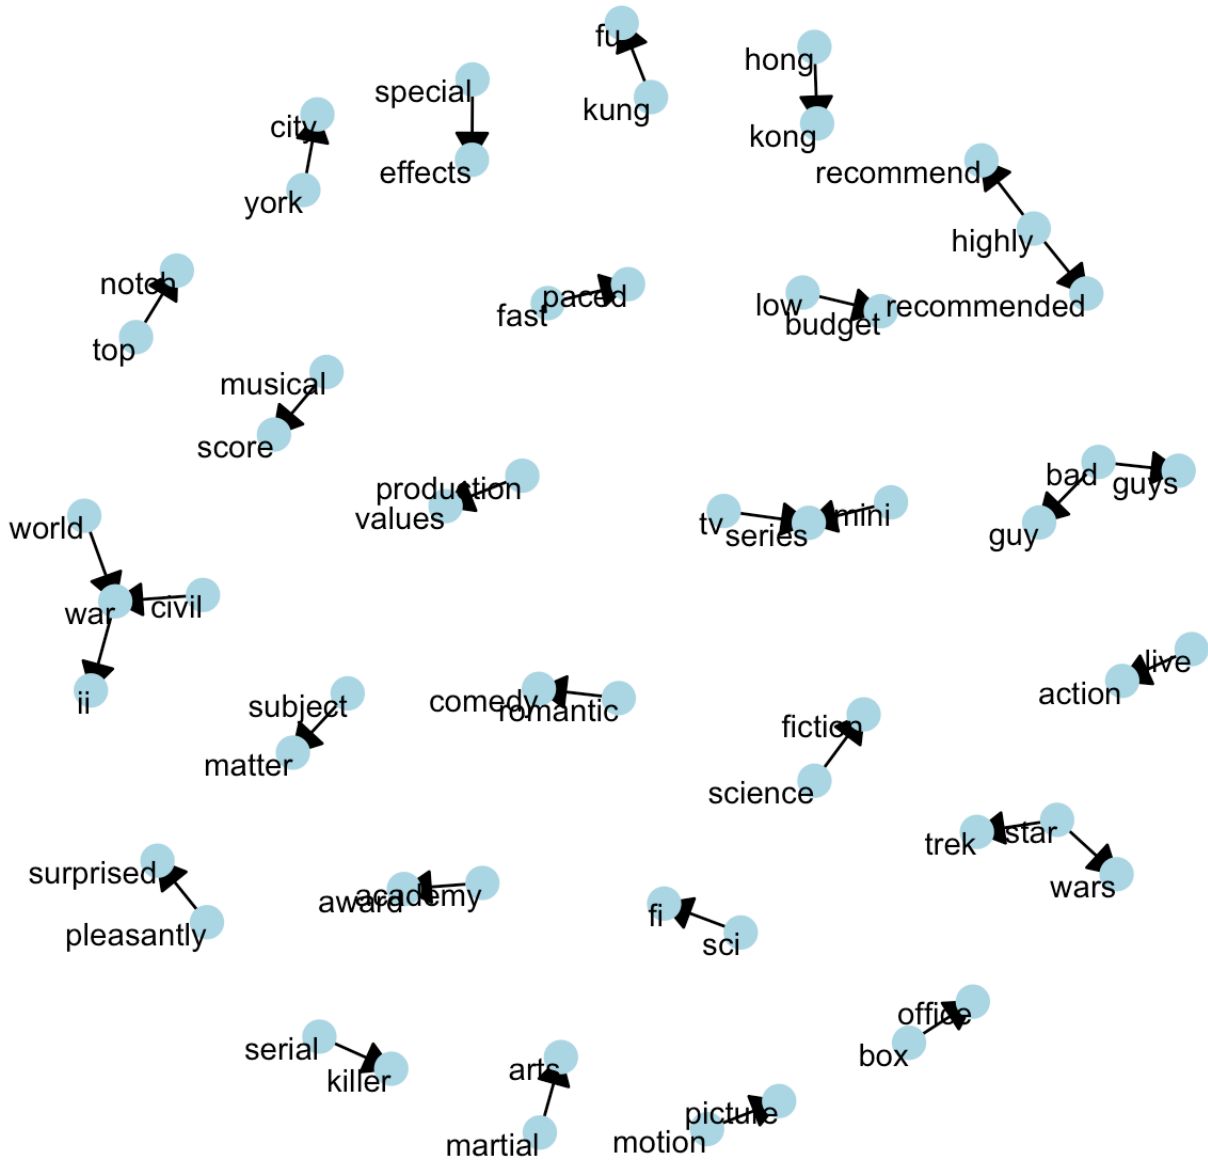
```
## # A tibble: 820,351 x 3
##   word1      word2      n
##   <chr>    <chr>   <int>
## 1 special effects  2240
## 2 low      budget  1812
## 3 sci      fi      1384
## 4 bad      acting   657
## 5 bad      guys     656
## 6 tv       series   587
## 7 bad      guy      573
## 8 production values  562
## 9 highly   recommend 553
## 10 world   war       529
## # ... with 820,341 more rows
```

bigram and visualisation for negative, positive remarks

```
## # A tibble: 465,153 x 3
##   word1    word2      n
##   <chr>   <chr>   <int>
## 1 special effects    804
## 2 sci      fi        620
## 3 low      budget    592
## 4 highly   recommend  502
## 5 world    war        380
## 6 highly   recommended 372
## 7 tv       series     337
## 8 top      notch      295
## 9 kung     fu        271
## 10 science fiction    257
## # ... with 465,143 more rows
```

```
## Warning: Using the `size` aesthetic in this geom was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` in the `default_aes` field and elsewhere instead.
```

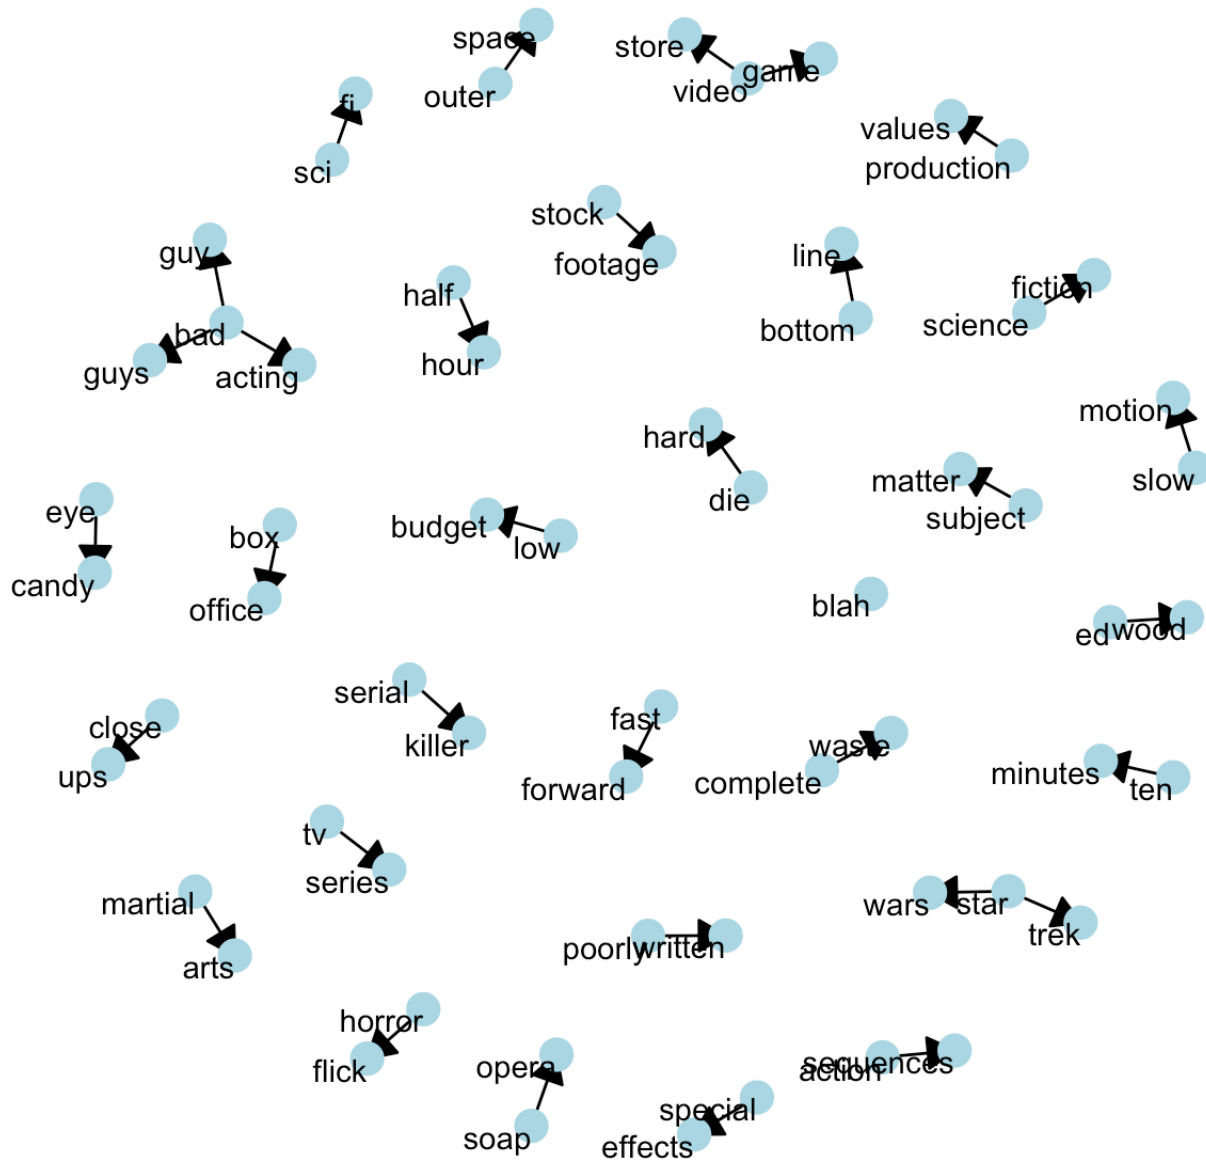
top bigrams (Frequency >= 150) in positive remarks



top bigram graph in negative remarks

```
## # A tibble: 416,404 x 3
##   word1      word2      n
##   <chr>    <chr>  <int>
## 1 special effects 1436
## 2 low      budget 1220
## 3 sci      fi      764
## 4 bad      acting 603
## 5 bad      guys   408
## 6 bad      guy    364
## 7 production values 345
## 8 90       minutes 308
## 9 20       minutes 290
## 10 martial arts    260
## # ... with 416,394 more rows
```

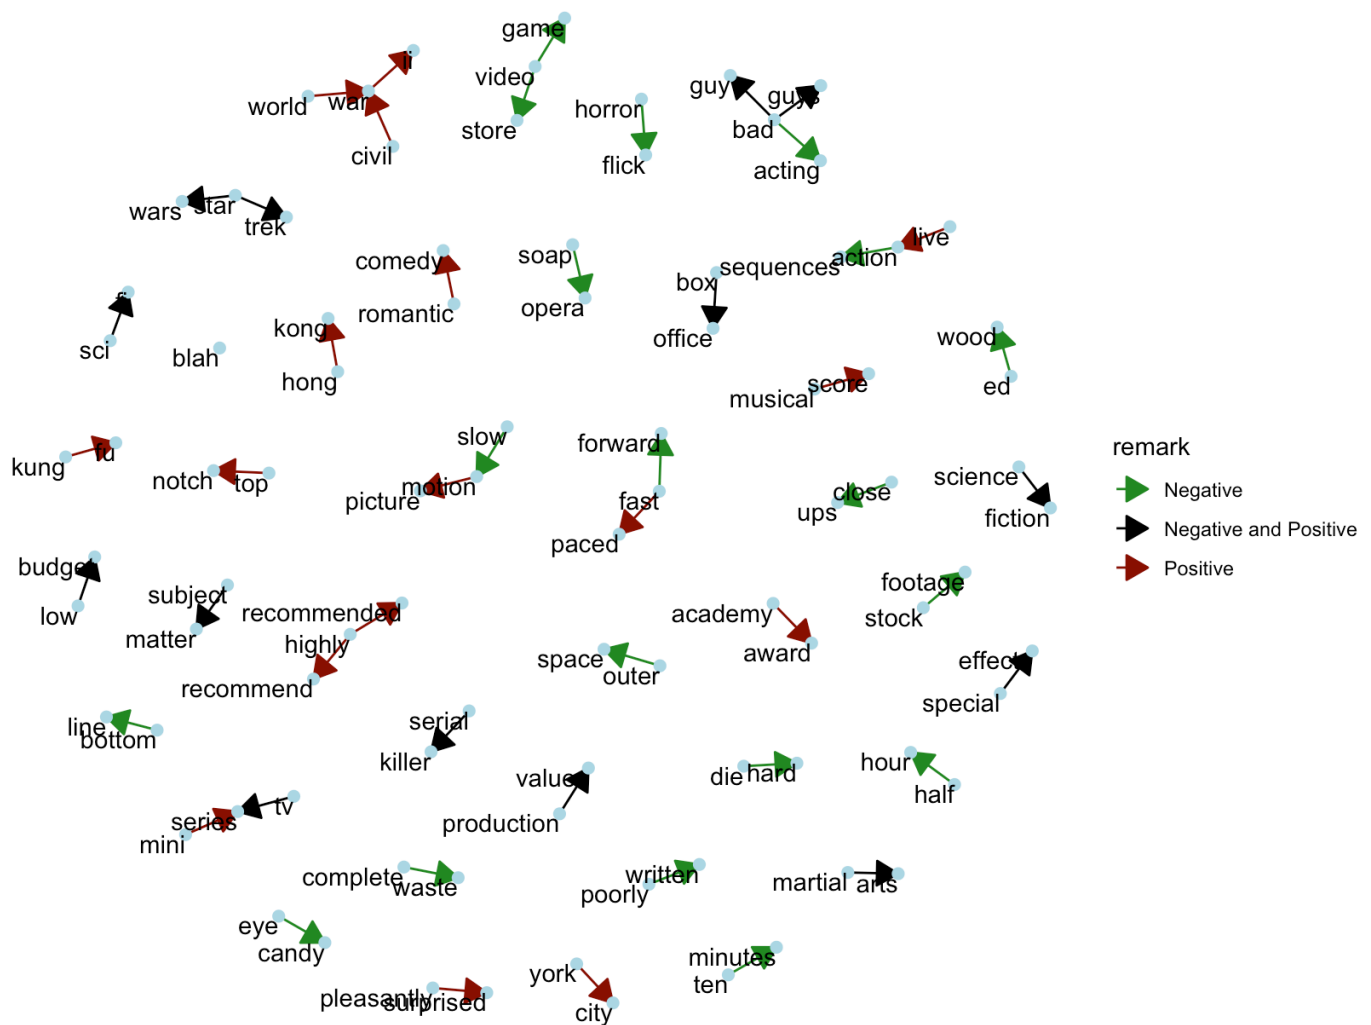

top bigrams (Frequency>=150) in negative remarks



From above two tables and graphs, we found some bigrams present both in negative remarks and positive remarks, for example, bigram of `special` and `effects` occupied the first place in both ranklist!

So we're about to show top bigrams in positive remarks only and negative remarks only and both in positive and negative remarks as following:

```
## # A tibble: 9 x 4
## # Groups:   remark [3]
##   word1    word2          n remark
##   <chr>   <chr>        <int> <chr>
## 1 bad      acting          603 Negative
## 2 fast     forward          247 Negative
## 3 ten      minutes          221 Negative
## 4 special  effects         2240 Negative and Positive
## 5 sci      fi             1384 Negative and Positive
## 6 low      budget          1812 Negative and Positive
## 7 highly   recommend         502 Positive
## 8 world    war              380 Positive
## 9 highly   recommended       372 Positive
```



In other analyses, we may want to work with the recombined words. tidyr's `unite()` function is the inverse of `separate()`, and lets us recombine the columns into one. Thus, "separate/filter/count/unite" let us find the most common bigrams not containing stop-words.

```
bigrams_united <- bigrams_filtered %>%
  unite(bigram, word1, word2, sep = " ")

bigrams_united
```

```
## # A tibble: 1,319,960 x 3
##   sentiment review_number bigram
##   <chr>          <int> <chr>
## 1 positive              1 faint hearted
## 2 positive              1 drugs sex
## 3 positive              1 oswald maximum
## 4 positive              1 maximum security
## 5 positive              1 emerald city
## 6 positive              1 experimental section
## 7 positive              1 glass fronts
## 8 positive              1 aryans muslims
## 9 positive              1 muslims gangstas
## 10 positive             1 gangstas latinos
## # ... with 1,319,950 more rows
```

In other analyses, we may be interested in the most common trigrams, which are consecutive sequences of 3 words. We can find this by setting `n = 3`:

```
IMDB_df %>%
  unnest_tokens(trigram, review, token = "ngrams", n = 3) %>%
  filter(!is.na(trigram)) %>%
  separate(trigram, c("word1", "word2", "word3"), sep = " ") %>%
  filter(!word1 %in% stop_words$word,
         !word2 %in% stop_words$word,
         !word3 %in% stop_words$word) %>%
  count(word1, word2, word3, sort = TRUE)
```

```
## # A tibble: 411,042 x 4
##   word1    word2    word3      n
##   <chr>   <chr>   <chr>   <int>
## 1 world   war       ii       230
## 2 sci     fi        channel  164
## 3 low     budget    horror    134
## 4 bad     acting    bad       103
## 5 mystery science theater    97
## 6 blah    blah     blah      90
## 7 texas   chainsaw massacre    90
## 8 blair   witch     project    84
## 9 local   video     store      80
## 10 tour    de        force      77
## # ... with 411,032 more rows
```

5. Converting to and from non-tidy formats

Just as some existing text mining packages provide document-term matrices as sample data or output, some algorithms expect such matrices as input. Therefore, tidytext provides `cast_verbs` for converting from a tidy form to these matrices.

```
IMDB_dtm <- tidy_IMDB %>%  
  count(review_number, word) %>%  
  cast_dtm(review_number, word, n)  
  
IMDB_dtm
```

```
## <<DocumentTermMatrix (documents: 50000, terms: 118916)>>  
## Non-/sparse entries: 3423798/5942376202  
## Sparsity           : 100%  
## Maximal term length: 72  
## Weighting           : term frequency (tf)
```

6. Topic modeling

Latent Dirichlet allocation is one of the most common algorithms for topic modeling. Without diving into the math behind the model, we can understand it as being guided by two principles.

This function returns an object containing the full details of the model fit, such as how words are associated with topics and how topics are associated with documents.

```
# choose number of topics, 6  
k = 6  
  
# set a seed so that the output of the model is predictable  
IMDB_lda <- LDA(IMDB_dtm, k, control = list(seed = 1234))  
IMDB_lda
```

```
## A LDA_VEM topic model with 6 topics.
```

```
#> A LDA_VEM topic model with 6 topics.
```

To start to exploring the `lda` model, we list the most frequent 30 terms in the topic listed, in rank order.

##	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6
## 1	acting	love	bad	bad	love	acting
## 2	life	feel	life	life	bad	series
## 3	funny	funny	acting	comedy	life	feel
## 4	music	minutes	horror	acting	set	original
## 5	entertaining	found	star	funny	role	role
## 6	times	bad	takes	minutes	world	bit
## 7	series	world	terrible	worst	music	guy
## 8	played	horror	action	guy	original	played
## 9	comedy	acting	play	script	script	tv
## 10	performance	original	guy	bit	start	music
## 11	hard	death	father	day	woman	fun
## 12	book	evil	american	hollywood	idea	times
## 13	bit	performance	piece	terrible	dvd	worth
## 14	shot	performances	lines	times	true	war
## 15	fan	worst	world	idea	half	bad
## 16	screen	plays	day	performance	action	book
## 17	plays	main	dead	read	mind	home
## 18	left	nice	audience	boy	family	read
## 19	job	family	head	shot	worst	low
## 20	completely	girl	excellent	reason	times	excellent
## 21	house	hard	rest	nice	found	true
## 22	script	sense	black	audience	horror	version
## 23	actor	guy	tv	sense	funny	girl
## 24	girl	classic	job	book	american	main
## 25	family	job	school	fine	version	mother
## 26	john	enjoy	funny	sex	live	performance
## 27	beautiful	series	money	effects	wrong	screen
## 28	watched	recommend	past	stupid	stuff	remember
## 29	mind	american	dialogue	hard	death	school
## 30	fun	life	plays	half	actor	idea

The tidytext package provides this method for extracting the per-topic-per-word probabilities, called (“beta”), from the model.

```
IMDB_topics <- tidy(IMDB_lda, matrix = "beta")
IMDB_topics %>% head()
```

```
## # A tibble: 6 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1 accustomed 0.0000102
## 2     2 accustomed 0.0000173
## 3     3 accustomed 0.00000972
## 4     4 accustomed 0.0000175
## 5     5 accustomed 0.0000144
## 6     6 accustomed 0.0000184
```

We could use `dplyr`'s `slice_max()` to find the 10 terms that are most common within each topic and create a visualization.



Now we use data after stemming and lemmatization to see if we have an improved results for topic identification. Because stemming would convert “plays” to “plai” which can be confusing for us to interpret the modeling results, we use lemmatized data to fit the model.

```
# LDA model fitting after stemming and lemmatization

lemma_dtm <- lemmatization%>% count(review_number, lemma) %>%
  cast_dtm(review_number, lemma, n)
# lemma_dtm still has sparsity of 100%, but has less terms: 100852

lemma_lda <- LDA(lemma_dtm, k, control = list(seed = 1234))
lemma_terms <- as.data.frame(topicmodels::terms(lemma_lda, 30),
                             stringsAsFactors = FALSE)

lemma_terms[1:k]
```

##	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6
## 1	funny	bad	bad	act	bad	play
## 2	love	love	star	bad	performance	bad
## 3	life	guy	life	book	screen	act
## 4	play	day	funny	life	kid	love
## 5	lead	woman	love	write	run	life
## 6	minute	play	start	set	life	kid
## 7	role	role	head	play	moment	reason
## 8	take	minute	feel	performance	talk	guy
## 9	leave	feel	hard	world	leave	run
## 10	act	job	line	kill	act	line
## 11	original	action	live	version	begin	girl
## 12	fan	find	take	eye	girl	laugh
## 13	feel	boy	action	funny	time	write
## 14	star	child	series	fun	set	comedy
## 15	woman	act	shoot	guy	horror	kill
## 16	start	fan	main	enjoy	hard	horror
## 17	audience	girl	day	time	music	mind
## 18	line	music	horror	leave	friend	couple
## 19	script	leave	short	series	world	talk
## 20	boy	fun	sense	bite	family	fall
## 21	shoot	comedy	play	remember	sound	find
## 22	actor	performance	picture	feel	love	enjoy
## 23	wrong	friend	tv	hope	completely	screen
## 24	happen	bring	write	house	laugh	call
## 25	run	time	family	death	low	money
## 26	kill	lead	classic	woman	version	true
## 27	call	original	piece	role	direct	bite
## 28	die	happen	low	wife	death	sense
## 29	black	script	dead	live	tell	idea
## 30	house	experience	original	comedy	series	view

```

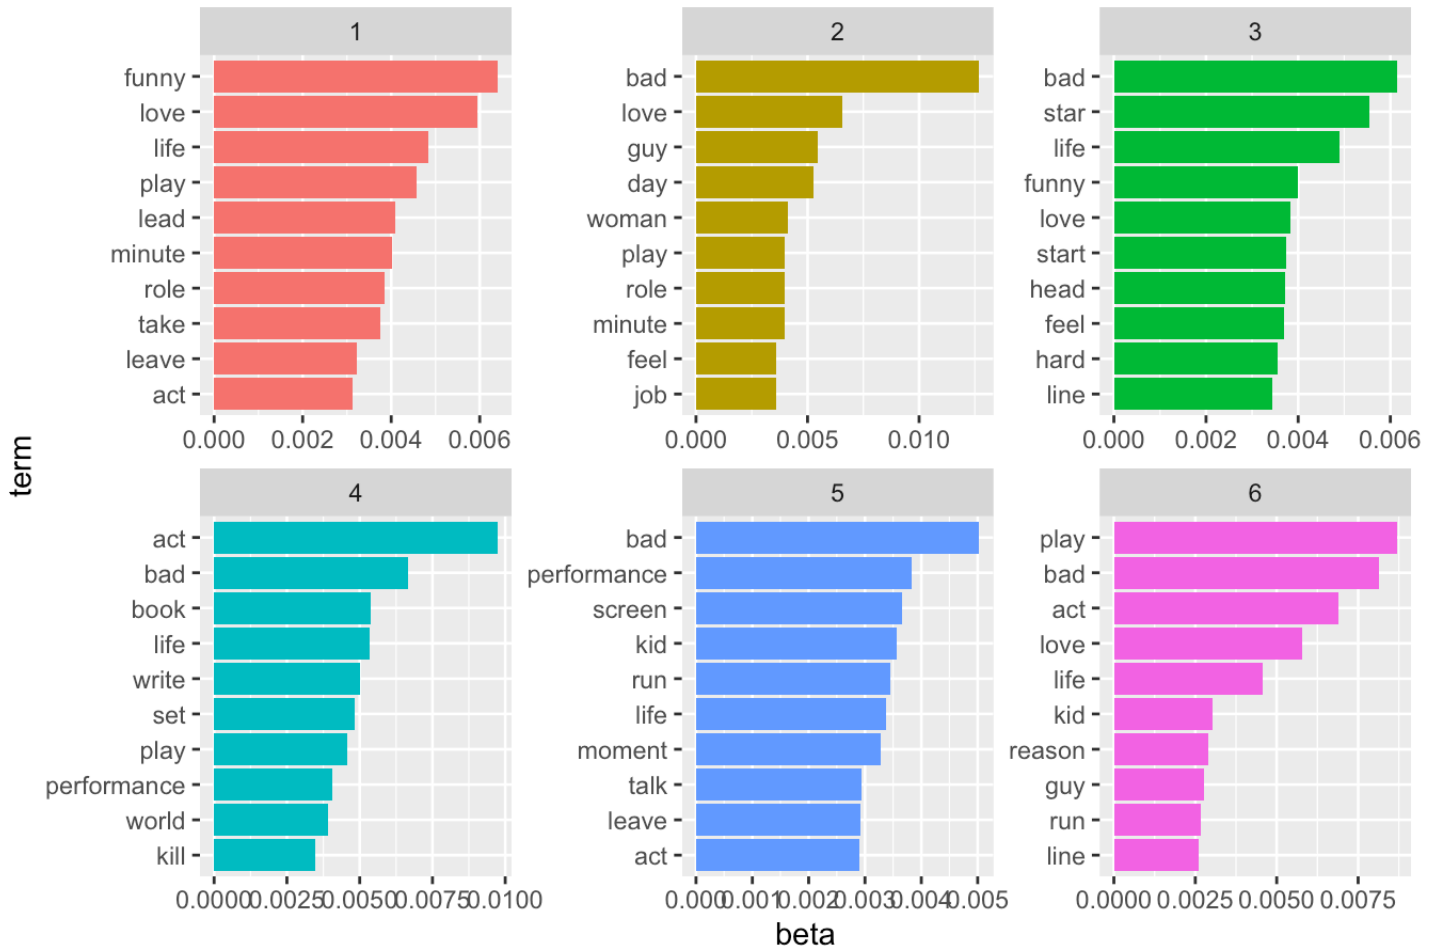
lemma_topics <- tidy(lemma_lda, matrix = "beta")

lemma_top_terms <- lemma_topics %>%
  group_by(topic) %>%
  slice_max(beta, n = 10) %>%
  ungroup() %>%
  arrange(topic, -beta)

lemma_top_terms %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
  ggplot(aes(beta, term, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  scale_y_reordered()+
  labs(title = "K=6, After lemmatization LDA model results")

```


K=6, After lemmatization LDA model results



The most common words for each topic now seems more exclusive and words like “played” don’t appear in the plot. However, based on the output, setting topic(K) to 6 cannot really help us distinguish one topic from the other one by eyeballing. So we increase K to 10 given the data size of our dataset and the results improve by having less common words associated with each topic.

##	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6
## 1	acting	love	bad	bad	bad	acting
## 2	music	minutes	acting	acting	love	bit
## 3	funny	funny	horror	life	world	music
## 4	bit	world	father	minutes	music	role
## 5	times	acting	life	comedy	life	played
## 6	played	feel	star	worst	role	series
## 7	life	horror	american	funny	set	tv
## 8	performance	worst	world	bit	original	original
## 9	screen	original	takes	script	true	worth
## 10	book	performance	play	guy	dvd	times
## 11	actor	sense	guy	day	script	true
## 12	series	bad	dead	times	worst	guy
## 13	enjoy	family	terrible	sense	family	feel
## 14	comedy	enjoy	money	read	live	war
## 15	shot	evil	sound	performance	times	fun
## 16	watched	main	dialogue	sound	actor	book
## 17	family	girl	action	shot	half	screen

## 18	father	plays	day	worth	mind	read
## 19	hard	actor	person	true	woman	version
## 20	girl	found	head	stupid	version	performance
## 21	finally	american	screen	hollywood	sense	excellent
## 22	beautiful	wife	picture	book	money	girl
## 23	script	hard	tv	short	match	low
## 24	friend	performances	past	reason	horror	home
## 25	plays	killer	excellent	tv	night	main
## 26	left	production	audience	terrible	finally	short
## 27	mind	guy	hope	effects	american	picture
## 28	job	series	loved	special	funny	family
## 29	version	fun	lines	entire	5	dead
## 30	fun	death	funny	person	awful	bad
##	Topic 7	Topic 8	Topic 9	Topic 10		
## 1	bad	series	bad	love		
## 2	funny	life	love	life		
## 3	life	acting	life	acting		
## 4	set	bad	comedy	bad		
## 5	horror	role	guy	action		
## 6	day	action	funny	horror		
## 7	guy	original	found	feel		
## 8	fan	shot	war	book		
## 9	role	screen	day	script		
## 10	family	feel	classic	original		
## 11	house	takes	fan	fun		
## 12	budget	woman	plays	times		
## 13	love	love	world	boy		
## 14	left	set	recommend	found		
## 15	completely	john	home	watched		
## 16	performance	found	entertaining	woman		
## 17	effects	classic	beautiful	black		
## 18	main	played	hard	death		
## 19	terrible	worst	fun	hard		
## 20	idea	idea	wonderful	play		
## 21	wrong	dvd	effects	john		
## 22	start	star	human	left		
## 23	series	reason	worst	idea		
## 24	girl	tv	acting	sense		
## 25	head	guy	times	low		
## 26	friends	girl	takes	main		
## 27	friend	american	tv	music		
## 28	game	beautiful	nice	job		
## 29	play	start	remember	short		
## 30	tv	house	played	completely		

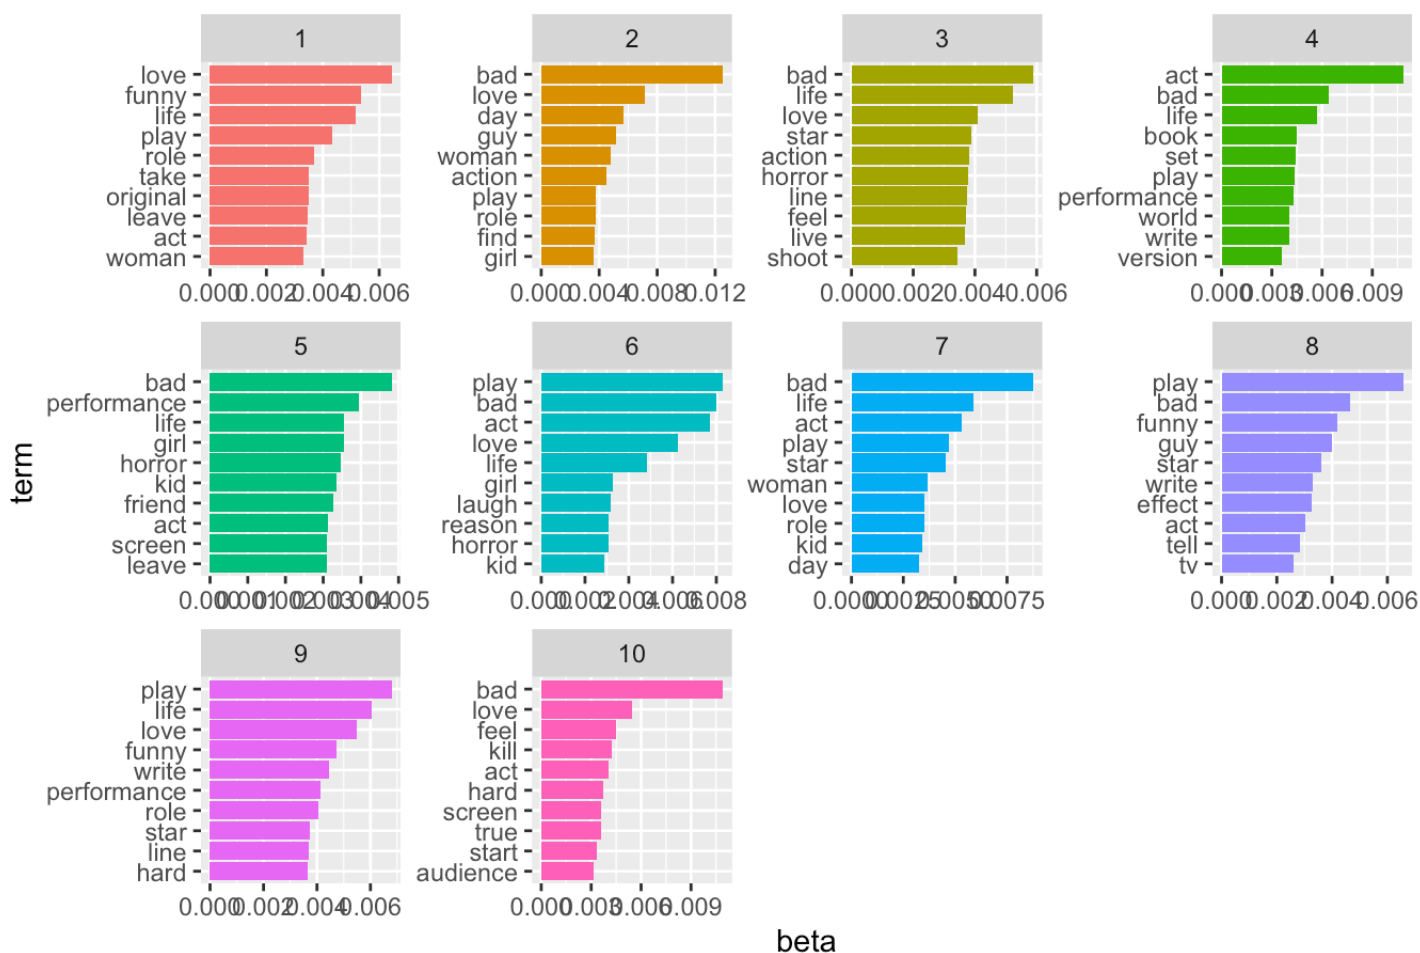
##	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5	Topic 6	Topic 7
## 1	love	bad	bad	act	bad	play	bad
## 2	funny	love	life	bad	performance	bad	life
## 3	life	day	love	life	life	act	act
## 4	play	guy	star	book	girl	love	play

## 5	role	woman	action	set	horror	life	star
## 6	take	action	horror	play	kid	girl	woman
## 7	original	play	line	performance	friend	laugh	love
## 8	leave	role	feel	world	act	reason	role
## 9	act	find	live	write	screen	horror	kid
## 10	woman	girl	shoot	version	leave	kid	day
## 11	script	feel	main	fun	laugh	line	job
## 12	line	fun	start	kill	run	comedy	music
## 13	happen	friend	funny	bite	version	mind	effect
## 14	feel	act	series	time	begin	guy	find
## 15	minute	minute	take	leave	time	couple	leave
## 16	shoot	happen	day	eye	family	money	time
## 17	fan	leave	tv	enjoy	love	kill	comedy
## 18	start	comedy	head	series	world	find	enjoy
## 19	die	boy	picture	hope	moment	call	lead
## 20	wrong	horror	short	feel	talk	fall	moment
## 21	lead	fan	family	woman	set	run	half
## 22	black	performance	play	guy	move	idea	guy
## 23	call	script	original	funny	music	bite	dvd
## 24	boy	music	sense	live	mind	enjoy	fall
## 25	expect	job	black	comedy	series	write	run
## 26	money	time	writer	excellent	low	script	death
## 27	laugh	original	war	remember	hour	friend	set
## 28	kill	move	classic	lose	script	talk	nice
## 29	actor	expect	hear	laugh	war	human	action
## 30	star	bring	low	tv	action	viewer	main

##	Topic 8	Topic 9	Topic 10
## 1	play	play	bad
## 2	bad	life	love
## 3	funny	love	feel
## 4	guy	funny	kill
## 5	star	write	act
## 6	write	performance	hard
## 7	effect	role	screen
## 8	act	star	true
## 9	tell	line	start
## 10	tv	hard	audience
## 11	leave	start	book
## 12	audience	bad	minute
## 13	night	feel	girl
## 14	take	lead	funny
## 15	world	kid	set
## 16	run	minute	music
## 17	boy	set	call
## 18	moment	begin	episode
## 19	series	world	family
## 20	sound	bite	run
## 21	fan	fan	war
## 22	classic	take	american
## 23	kill	head	friend

```
## 24      minute      live      comedy
## 25      family      black      fight
## 26 performance      guy      lead
## 27      laugh      short      direct
## 28      lead      act      child
## 29      edit      eye      video
## 30      version      actor      view
```

K=10, After lemmatization LDA model results



Besides estimating each topic as a mixture of words, LDA also models each document as a mixture of topics. We can examine the per-document-per-topic probabilities, called (“gamma”), with the matrix = “gamma” argument to tidy().

```
IMDB_gamma <- tidy(IMDB_lda, matrix = "gamma")
IMDB_gamma
```

```
## # A tibble: 500,000 x 3
##   document topic  gamma
##   <chr>      <int>  <dbl>
## 1 1 1          1 0.102
## 2 2          1 0.100
## 3 3          1 0.101
## 4 4          1 0.101
## 5 5          1 0.0988
## 6 6          1 0.0999
## 7 7          1 0.0981
## 8 8          1 0.100
## 9 9          1 0.101
## 10 10         1 0.0995
## # ... with 499,990 more rows
```

7. Topic Modeling extension: Determine k number of topics

In section 6, we assume the number of topics to be 6 to conduct the LDA since not pre-topic tags were provided. However, one major aspect of LDA, is that we need to know the exact k number of optimal topics for the documents. In order to accomplish this task, we are going to use a harmonic mean method to determine k based on Martin Ponweiser's thesis.(see README.md file)

First, we set up the fuction for computing the harmonic mean

```
# The harmonic mean function
harmonicMean <- function(logLikelihoods, precision = 2000L) {
  llMed <- median(logLikelihoods)
  as.double(llMed - log(mean(exp(-mpfr(logLikelihoods,
                                     prec = precision) + llMed)))))
}
```

In order to find the best value for k, we do this over a sequence of topic models with different vales for k. This will generate numerous topic models with different numbers of topics, creating a vector to hold the k values.

```

up_k <- 10

# We will use a sequence of numbers from 2 to up_k
seqk <- seq(2, up_k, 1)
burnin <- 1000
iter <- 1000
keep <- 50
fitted_many <- lapply(seqk, function(k) topicmodels::LDA(IMDB_dtm, k = k,
method = "Gibbs", control = list(burnin = burnin, iter = iter, keep = keep)))

# extract logliks from each topic
logLiks_many <- lapply(fitted_many, function(L) L@logLiks[-c(1:(burnin/keep))])


# compute harmonic means
hm_many <- sapply(logLiks_many, function(h) harmonicMean(h))

```

We could visualize the results of harmonic means by plotting the results

Latent Dirichlet Allocation Analysis of IMDB

How many distinct topics in the Reviews?



From the plot above, we can see that the optimal number of topics for our model is 10.

The following code returns the optimal number of topics:

```
seqk[which.max(hm_many)]
```

```
## [1] 10
```